

GPU-Based Parallel Kalman Filter

Zhiyuan Lin, David Moore, Stuart Russell

University of California, Berkeley

Abstract—We describe a parallel GPU Kalman filtering implementation using splash belief propagation to parallelize computation across timesteps, allowing a tradeoff between accuracy and speed. Focusing on a real world application to compute the marginal likelihood of seismic waveforms, we optimize the GPU implementation taking advantage of global memory access patterns and GPU fast memory. Our parallel Kalman filter can achieve nearly linear speedup in simple applications and can outperform a CPU implementation by an order of magnitude in some cases.

I. INTRODUCTION

Kalman filtering [1] is a well known algorithm that has been widely used in many applications for state estimation of linear systems. It uses a series of observations over time to produce estimates of unknown states. More precise than those based on a single observation alone. However, the matrix operations involved in Kalman filter updates can be a computational burden. In order to solve this problem, we explore potential ways to parallelize the algorithm.

A natural way to make a Kalman filter faster is to parallelize it on a Graphics Processor Unit (GPU). Most existing ways to parallelize the Kalman filter simply parallelize its matrix operations. Considering the cost of transferring data between GPU and CPU, parallelizing each matrix operation might not have performance gain for some applications, especially when the matrices are not large enough to occupy all GPU cores.

This paper proposes a new way to parallelize a Kalman Filter. Inspired by Parallel Splash Belief Propagation [2], we parallelize over time instead of parallelizing matrix operations at each time step. Kalman filtering is a message passing algorithm. In many applications, it is reasonable to assume that the current state estimate relies more on recent states, and that states from long ago don't contribute much to the current state. Based on this assumption and the abundance of GPU threads, we compute a filtered state estimate for each timestep in a separate thread, based on the most recent observations. We implement this parallel Kalman Filter and optimize the GPU implementation, focusing on an application to computing the marginal likelihood of seismic waveforms. We compare the speed and accuracy of our parallel Kalman Filter with optimized CPU implementations.

The remaining sections are organized as follows. In section 2, we introduce the necessary background of this paper, including a brief description of the Kalman Filter, and analyze existing ways of parallelizing similar algorithms. In section 3, we describe the basics of the application model we are working on. In section 4 we discuss the details of our implementation. Section 5 explores the ways to optimize our implementation on a GPU. Section 6 performs various benchmarks and evaluate

the results. Section 7 makes a conclusion and give directions for future work.

II. BACKGROUND

A. Kalman Filter

The Kalman filter algorithm involves two phases: predicting and updating, corresponding to the transition model and observation model respectively. A Kalman filter takes in a series of noisy observations to estimate hidden states. In the prediction step, it predicts the estimate of the current hidden state, based on the updated estimate of the previous state. Given the noisy observation at the current time step, this predicted estimate is updated using a weighted average of the existing estimate and the observation. The Kalman gain is the weight that measures how much the new observation needs to be taken into consider. We provide details in the predict phase and update phase. We use \mathbf{k} to demonstrate each time step, and we will use these notations defined in Algorithm 1 in the following part of this paper.

B. Approaches of parallelization(Related Works)

Existing approaches of parallelizing Kalman filter focus on parallelizing matrix operations. We explored parallelism in graphical model inference, and was enlightened by Parallel splash belief propagation.

1) *Parallelize matrix operations*: The most popular way of parallelizing a Kalman filter on GPU is to parallelize its matrix multiplication and matrix inversion in predict phase and update phase at each time step [3], [4]. But when matrix dimension is not large enough to fill all the GPU pipelines, it's hard to achieve optimal performance. Also, data transfer between CPU and GPU can be a bottleneck, for this way of parallelizing needs to transfer data at each time steps.

2) *Break data dependency*: An implementation of a Kalman filter for single output system on multicore computational platforms [5] breaks the data dependencies through re-organizing calculations, by which an almost completely parallel algorithm is obtained. However, this is only suitable for a specific kind of Kalman filter applications.

3) *Parallel hidden Markov model(Reduce operations)* : Hidden Markov model shares many characteristics of Kalman filter, thus its parallelism may give us some insights on parallelizing Kalman filter.

CuHMM [6] is a CUDA implementation of hidden Markov model (HMM) training and classification, which is also ends up parallelizing matrix operations and has great speedup.

Another paper [7] describes algorithms for a parallel implementation of hidden Markov models with a small state space. It shows how to format the HMM algorithm using linear algebra,

Algorithm 1 Serial Kalman Filter**notations for Kalman filter:**

(m - dimension of hidden state; n - dimension of observation)

F - state transition matrix (m*m)

H - observation matrix (n*m)

Q - covariance matrix of process noise (m*m)

R - covariance matrix of observation noise (n*n)

 Z_k - observation at time step k

k - for each time step

function KALMAN PREDICT ($\hat{X}_{k|k}$, $P_{k|k}$)**returns** predicted state estimate**variables:** $\hat{X}_{k|k-1}$, predicted state estimate $\hat{X}_{k|k}$, updated state estimate $P_{k|k-1}$, predicted state estimate covariance $P_{k|k}$, updated state estimate covariance**inputs:** $\hat{X}_{k|k-1}$, $P_{k|k-1}$

$$\hat{X}_{k|k-1} = F\hat{X}_{k-1|k-1}$$

$$P_{k|k-1} = FP_{k-1|k-1}F^T + Q$$

function KALMAN PREDICT ($\hat{X}_{k|k}$, $P_{k|k}$)**returns** updated state estimate, predicted observation estimate**variables:** $\hat{X}_{k|k-1}$, $\hat{X}_{k|k}$, $P_{k|k-1}$, $P_{k|k}$, \hat{Y}_k , predicted observation estimate S_k , predicted observation estimate covariance K_k , optimal Kalman gain**inputs:** $\hat{X}_{k|k-1}$, $P_{k|k-1}$, Z_k

$$\hat{Y}_k = H\hat{X}_{k|k-1}$$

$$S_k = HP_{k|k-1}H^T + R$$

$$K_k = P_{k|k-1}H^TS_k^{-1}$$

$$\hat{X}_{k|k} = \hat{X}_{k|k-1} + K_k(Z_k - \hat{Y}_k)$$

$$P_{k|k} = (I - K_kH)P_{k|k-1}$$

function SERIAL KALMAN FILTER ($\hat{X}_{k|k}$, $P_{k|k}$)**returns** marginal likelihood (sum)**variables:** $\hat{X}_{k|k-1}$, $\hat{X}_{k|k}$, $P_{k|k-1}$, $P_{k|k}$, \hat{Y}_k , S_k ,*likelihood*, accumulated marginal likelihood**inputs:** Z, observations

prior, the prior distribution on the initial state

n, total time steps (length of signal)

likelihood $\leftarrow 0$ **for** k = 0 to n-1 **do****if** k = 0 **then** $(\hat{X}_{k+1|k}, P_{k+1|k}) \leftarrow$ prior**else** $(\hat{X}_{k+1|k}, P_{k+1|k}) \leftarrow$ \leftarrow KALMANPREDICT ($\hat{X}_{k|k}$, $P_{k|k}$)**end if** $(\hat{X}_{k+1|k+1}, P_{k+1|k+1}, \hat{Y}_k, S_k)$ \leftarrow KALMANUPDATE ($\hat{X}_{k+1|k}$, $P_{k+1|k}$)*likelihood* \leftarrow *likelihood* $+ \text{MULTIVARIATE_NORMAL.LOGPDF}(Z_k, \hat{Y}_k, S_k)$ **end for****return** *likelihood***end function**

which naturally lends the algorithm itself to parallelization. It uses parallel reduction [8] on one algorithm in HMM.

C. Parallel splash belief propagation

Our parallel model is based on parallel splash belief propagation [2].

Considering the natural, synchronous parallelization of belief propagation is highly inefficient. Parallel splash belief propagation focuses on parallel graphical model inference. By bounding the achievable parallel performance in chain graphical models, they develop a theoretical understanding of the parallel limitations of belief propagation. Their assumptions is, for a long chain graphical model with weak edge potentials, distant vertices are approximately independent. For a particular vertex, an accurate approximation to the belief may often be achieved by running belief propagation on a small subgraph around that vertex, which can reduce the sequential component of belief propagation to the longest path in the subgraph.

This parallelism is abstract, without respect to a particular probability model. Kalman filter is an example of a “chain” graphical model. We develop our parallel model based on this parallel splash belief propagation idea.

III. PROGRAMMING MODEL**A. GPU and CUDA**

GPU-accelerated computing uses a GPU to perform computation traditionally handled by the CPU. It puts compute-intensive part of the application to a GPU, while the remainder of the code still runs on the CPU. A CPU is optimized for sequential serial processing while a GPU has massive parallel architecture consisting of thousands of smaller, more efficient cores to handle multiple tasks simultaneously. We target GPU because it is designed for massive parallel calculation.

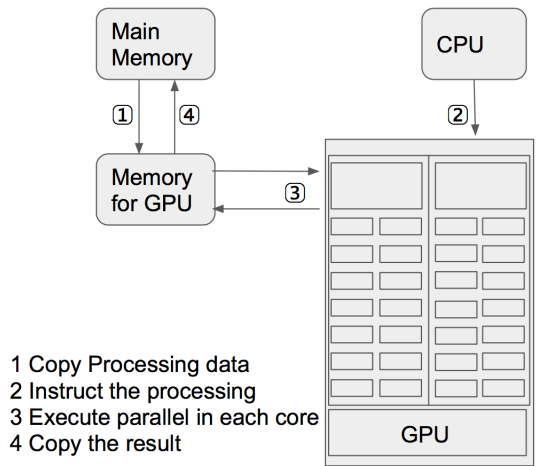


Figure 1: CUDA processing flow

CUDA [9] is a parallel computing platform and programming model invented by NVIDIA. It enables programmer to use GPU for general purpose processing. The CUDA platform can work with programming languages such as C,C++ and

Fortran, which makes GPU programming easier. We implement this parallel Kalman Filter model using Pycuda, which maps all of CUDA into Python.

B. Parallel Belief Propagation in Kalman Filter

This paper introduces a novel way to parallelize Kalman filter. We parallelize over time instead of parallelizing matrix operations at each time step. In many applications, the current state estimate relies more on recent states than on states from long ago. Based on this and considering a GPU has thousands of threads, we compute a filtered state estimate for each timestep in a separate thread, as is described in Algorithm 2. The work to compute one filtered state estimate is done by one single thread, based on the most recent observations. The number of observations required is the number of iterations of getting a result within some deviation.

This parallel approach is designed for Kalman filter applications that have some observations come at one time. After these observations arrive, the parallel Kalman filter do filtering for hidden states of these timesteps in parallel. In our approach, we are computing marginal likelihood for applications.

We can think that each thread is dealing with a small Kalman filter that compute the marginal likelihood for one hidden state, and we have many small Kalman filters running concurrently. Since we don't know the status of the initial state of these small Kalman Filters, each thread start out with uniform distribution with a zero mean and a covariance matrix with infinite variances. However, standard Kalman filter doesn't work with infinite variance beliefs, so we model this initial state into a Gaussian distribution with a zero mean and a covariance matrix with user specified large variances. Once we have this initial observation-based belief for each hidden state, passing message forward is just the standard Kalman Filter prediction step.

This parallelism does introduce some redundancy. Threads are doing overlapped computation, but we haven't found an elegant way to optimize it so far.

C. Accuracy

One important parameter in this parallel Kalman filter is the number of iterations of getting marginal likelihood result within certain accuracy. There is a tradeoff between accuracy and speed. Most applications will be more accurate and might converge to true value after some iterations. We can either derive how many iterations are needed to get true value from a specific model, or get it empirically after running experiments.

$$\text{Theoretical speedup} = \text{totalTimeStep} / \text{Iterations}$$

IV. OPTIMIZATION

In order to get optimal performance of GPU, the program should expose sufficient parallelism, use memory efficiently, minimize data transfer, and avoid different execution paths within the same warp. Coalescing global memory and minimizing redundant accesses to global memory are crucial.

Algorithm 2 Parallel Kalman Filter

function PARALLEL KALMAN FILTER (t, Z)

returns the marginal likelihood at timestep[t]

input:

t , the thread id for each cuda thread

Z , observations

persistent: F, H, Q, R ,

prior, a zero mean, and a covariance matrix with large variances

n , iteration of filtering in each thread

local variables (for each thread, and k is the timestep):

$\hat{Y}_k, S_k, K_k, \hat{X}_{k|k-1}, \hat{X}_{k|k}, P_{k|k-1}, P_{k|k}$,

logpdf, marginal likelihood for hidden state t

each cuda thread(t) do:

for $i = 0$ to $n-1$ **do**

if $i = 0$ **then**

$(\hat{X}_{t-n+1|t-n}, P_{t-n+1|t-n}) \leftarrow \text{prior}$

else

$(\hat{X}_{t-n+i+1|t-n+i}, P_{t-n+i+1|t-n+i})$

$\leftarrow \text{KALMANPREDICT}(\hat{X}_{t-n+i|t-n+i}, P_{t-n+i|t-n+i})$

end if

$(\hat{X}_{t-n+i+1|t-n+i+1}, P_{t-n+i+1|t-n+i+1}, \hat{Y}_k, S_k)$

$\leftarrow \text{KALMANUPDATE}(\hat{X}_{t-n+i+1|t-n+i}, P_{t-n+i+1|t-n+i})$

end for

return MULTIVARIATE_NORMAL.LOGPDF(Z_t, \hat{Y}_k, S_k)

end function

A. Expose Sufficient Parallelism

GPU is a parallel machine. It has lots of arithmetic pipelines. In order to exploit the abundance of pipelines, code must expose sufficient parallelism to occupy all pipelines and hide latency of pipelines. In our seismic signal model, we hope to run 15000 threads in parallel, which exposes GPU with sufficient parallelism. On the contrary, if we parallelize Kalman filter by parallelizing matrix operations, the matrices are not large enough to fill GPU threads.

B. Minimize data transfer (between host and device)

Usually, the first step to parallelize something on GPU is to identify the hotspots on serial CPU code, and parallelize them. Communication between host (CPU) and device (GPU) is expensive. Take this fact into consideration, we put all computation tasks in our model on the GPU. In this case, the CPU only needs to communicate with the GPU when all the work have done. Besides the constant transition and observation models, host only sends observations to device, and then host gets back the marginal likelihood of each timestep after device finish filtering.

C. Optimize Launch Configuration

Key of CUDA execution model is the difference between thread, threads block, grid and warp. Thread is a sequential execution unit, while threads block is a group of threads, and a grid is a collection of thread blocks. Threads within a block execute on a single Streaming Multiprocessor and

can communicate and cooperate with each other, while thread blocks do not synchronize with each other, thus threads in different thread blocks can not communicate with each other. A warp is a group of 32 threads and a thread block consists of several warps. Another essential part of CUDA execution model is the kernel, which represents the function that runs on GPU. When running a CUDA program, a CUDA kernel is launched as a grid of thread blocks.

Threads don't run simultaneously on GPU. On a multiprocessor, only a warp is executed physically in parallel (SIMD). GPU architecture hides latency from other warps on this multiprocessor to achieve high throughput. That is, when some warps are waiting for data, GPU executes the warp which is ready.

When configure kernel launch, thread block size should be a multiple of warp size. We choose block size based on the experiment result. 128 or 256 threads per block usually have a good performance in our model. And the grid size is total number of thread blocks.

D. Coalescing Global Memory Access

GPU Global memory capacity is around several GB, and is accessible by all threads, though has high latency. Operations on device are issued per warp. First, threads in a warp provide memory addresses, then the hardware determines which lines of memory are needed. Scattered address patterns or patterns with large strides between threads should be avoided so that a warp can access within a contiguous region. Performance gain can be achieved by coalescing global memory access [10] within a warp, the pattern is shown in Figure 2.


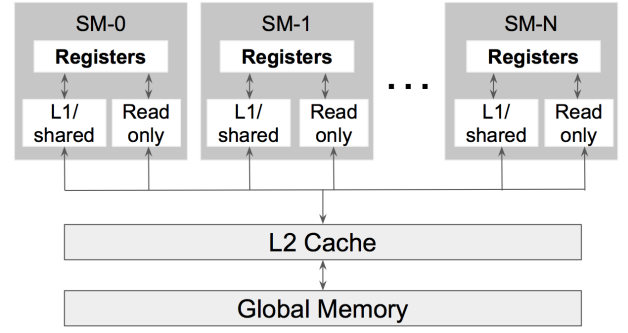


Figure 2: Global Memory Access Pattern

We have tried three memory allocation patterns to allocate matrices and arrays in our model: allocate in thread local memory; taking advantages of coalesced global memory as described above; and using standard global memory.

In our first implementation, we allocated most matrices and arrays in thread local memory. When a thread creates an array in CUDA kernel, this array is stored in thread local memory. Though theoretically local memory is part of global memory, the so called “global memory” is allocated outside kernel and can be accessed by all threads, while local memory is allocated within a single thread and can only be accessed by this specific thread.



(Streaming Multiprocessor - SM)

Figure 3: GPU Memory Hierarchy

Then we realized that using coalesced global memory would take much less time, and we changed the way we allocate memory. This modification did lead to better performance. However, we found that allocating some arrays from thread local memory to coalesced global memory achieved great performance gain, while others didn't.

In order to figure out why a better memory allocation pattern does not necessary improve performance, we have tried allocating global memory not in a coalesced pattern. It turned out that this patterned was twice as slow as allocating thread local memory. Our understanding is, though theoretically, latency of accessing local memory should be the same as global memory, registers are partitioned among threads, whose storage are local to threads. Some of the local arrays might be stored in registers, so local memory is faster than normal global memory allocation pattern. When changing the un-coalesced memory to coalesced memory, on average, it will be twice as fast as before.

The different performances between different arrays in our model can be explained as the different ways of using these arrays in kernel code. Considering this, we optimized memory allocation pattern empirically.

E. Exploit GPU Fast Memory

1) *GPU Memory Hierarchy* (Figure 3): We explore a typical NVidia GPU memory hierarchy to get an idea of how fast each memory is. Register is the fastest memory on GPU, but it's managed by compiler, and its storage is local to each thread. Shared memory and L1 cache also have low latency (e.g., 1-2cycles) and high bandwidth (e.g., 2.5TB/s) and it is program-managed. L2 cache is slower than memory mentioned above, and it's hardware-managed. Global memory is large, accessible by all threads and CPU, but has higher latency (e.g., 400-800 cycles) and lower bandwidth (e.g., 250GB/s). GPU caches are not intended for the same use as CPU caches, because it is much smaller, especially for each thread, and we will explore this more below.

2) *Shared Memory*: Shared memory is the fastest configurable GPU memory, but it's small. In our case, it can only accommodate several single value variables. We have tried to put these eligible variables into shared memory. However,

in all cases, it only became 1% faster, which is the same as configuring the cache to “prefer L1”. Considering using shared memory will have more complex configuration, we simply set the flag to give more cached memory to L1 cache, and let compiler to make its own decision.

3) *Constant Memory*: NVIDIA hardware provides 64KB of constant memory. Though constant memory resides on global memory, it is cached, which is different from standard global memory. Reading from constant memory is issued per half warp. If threads in a half warp want to read the same memory address, reading from constant cache is as fast as reading from registers. On the contrary, if threads in a half warp access different addresses, operations will be serialized, which is much slower. In order to maximize performance, we put read-only variables in constant memory if possible, which make the performance become 20% faster.

V. EVALUATION MODEL

In this section, we introduce concrete models that we have used to evaluate the parallel Kalman filter.

A. Seismic Signal Model

We model a seismic signal at a single station as a background noise process, plus arriving seismic phases. In this paper, we implement parallel Kalman filter based on this model. This probability model constructs the marginal distribution on the noisy signal as an explicit Gaussian, and compute the required likelihoods and posterior distributions using a Kalman filter. The background noise is modeled into an autoregressive process. We will briefly describe the AR process in next subsection. For more information about this model, please see this pending paper “Linear Gaussian Signal Model”, which describes this seismic signal model in details.

We have implemented the GPU-based parallel Kalman filter based on this seismic model using Pycuda. We reduced the float operations in the Seismic signal model by customizing matrix structure and simplifying observation model.

B. Autoregressive Process

An autoregressive process [11] of order p is a stochastic process with no hidden state, in which the expected value at time t is a linear function of the values at times $t-p, \dots, t-1$, with Gaussian noise at each step.

$$z_t = \sum_{k=1}^p \phi_k z_{t-k} + \varepsilon_t, \\ \varepsilon_t \sim N(0, \sigma_n^2).$$

C. Random Walk

Random walk [12] is a stochastic activity that consists of a succession of random steps. When model random walk into Kaman filter, its state transition matrix is identity matrix, and the observation is the hidden state at that time step plus some noise.

VI. EVALUATION

A. Setup

For parallel evaluations, we ran benchmarks on a server with a NVIDIA Quadro K4200 GPU, which has 1344 GPU cores, 4GB GPU memory with 173GB/s bandwidth and 64KB of RAM (configurable partitioning of shared memory and L1 cache), and with an Intel Xeon Processor E3-1271 v3.

For all the serial evaluations, we ran benchmarks on a computer with 2.6GHz Intel Core i5 Processor and 8GB 1600MHz memory.

B. Performance - Speed




Figure 4: Performance of Random Walk

1) *Toy example - Random Walk*: We first compare the performance of serial random walk (RW) and parallel RW on both 2 dimensional and 10 dimensional RW. The number of iterations is 10 here, at which the marginal likelihood converges to true value.

As shown in Figure 5, the serial python implementation has approximately linear complexity, while cuda implementation is much faster and scalable. Processing a 4096 steps two dimensional RW in parallel only takes 0.05X longer than processing a 32 steps signal in parallel. For ten dimensional RW, it achieves 470X speedup, running in 4096 threads taking 2.8X as long as running in 32 threads.

The speedup in parallel RW is great and sometimes the actual speedup can even exceed theoretical speedup. One factor that could account for this is that Python has interpreter overhead, thus it's much slower than CUDA implementation. Despite of this, we can still see great advantages when running the parallel Kalman filter.

2) *Autoregressive Process*: AR process is an essential part of the seismic model. We set the AR process to 10 order. In order to achieve accuracy, according to Figure 6, we choose the number of iterations as 11 for parallel version.

We tested our GPU-accelerated Kalman Filter against a fully optimized CPU-based Kalman Filter. One trick the C++ CPU code does is that it stops computing covariance after the




Figure 5: Performance of Autoregressive Process

covariance matrix converges. If there is a missing observation, it will re-compute the covariance matrix until it converges again, losing performance. We compared the performance of parallel version and serial version in both the fully observed setting and with missing observations every 20 time steps.

As we can see in the graph, the parallel version outperforms the serial version greatly. It can achieve up to 7X speedup when there are some missing observations. As we expect, the parallel version achieves higher speedup when observations are not fully observed. There is a tradeoff between speed and accuracy: here we set number of iterations to 11; one can get a higher speedup by iterating less, which may have less accurate result. We will explore more about accuracy in the following part.

C. Correctness

Tradeoff between accuracy and speed is crucial in our parallel Kalman filter, while it makes this parallel Kalman filter more flexible. From Figure 6, we can see in the 168 dimensional seismic signal, 10-order autoregressive process, and 2 dimensional random walk, when numbers of iterations increase, the marginal likelihood absolute deviations decrease. The marginal likelihood converges after a few time steps. From all the experiments we have ran, we can deduce that this GPU-

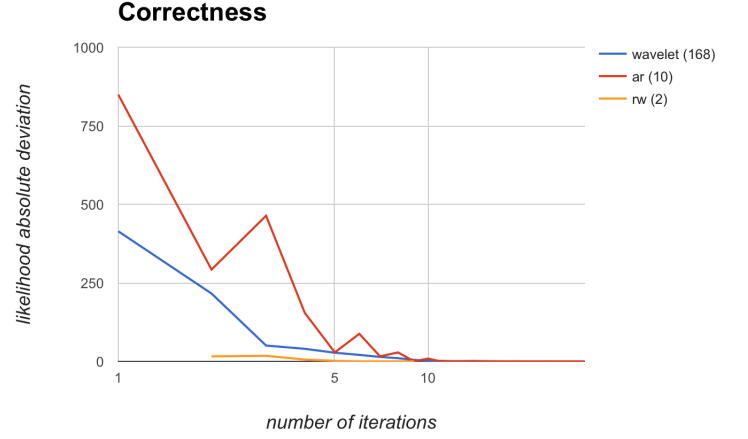


Figure 6: Correctness

based parallel Kalman filter has higher performance gain when dealing with applications that can converge in less time steps.

VII. CONCLUSION AND FUTURE WORK

We have presented a new approach to parallelize Kalman filter in this paper, leveraging general purpose GPU. We parallelize it over time instead of parallelizing matrix operations. While this parallel approach is only designed for applications that have some observations come at one time, we have found it is powerful enough to express these Kalman Filter applications.

We have evaluated our model by implementing it in Pycuda and tested it against optimized CPU code. It's shown that our model has better overall performance in small dimension and long signal length applications. Also, tradeoff between correctness and performance is interesting in our algorithm.

We have already seen desirable performance gain of our parallel Kalman filter, and here are several interesting directions for future research.

1. Parallelize matrix operations when the matrix is large, and combine it with our existing parallel model. To be more specific, in a long signal, parallelize the matrix operations in time steps that have large matrices, which corresponds to arriving earthquake signal in the seismic signal, and parallelize over time when matrices are small, which corresponds to the phase that only has background noise in our application.
2. Generalize the approach to fit more applications.
3. Our parallel model does some redundant computation in each thread, which takes lots of time. We should find a way to cut the redundancy in the future.

REFERENCES

- [1] S. Russell and P. Norvig, "Kalman filter," *Artificial intelligence: a modern approach*, 3rd ed, pp. 584–590, 1995.
- [2] J. Gonzalez, Y. Low, and C. Guestrin, "Parallel splash belief propagation," DTIC Document, Tech. Rep., 2010.
- [3] T. Blattner and S. Yang, "Performance study on cuda gpus for parallelizing the local ensemble transformed kalman filter algorithm," *Concurrency and Computation: Practice and Experience*, vol. 24, no. 2, pp. 167–177, 2012.

- [4] M.-Y. Huang, S.-C. Wei, B. Huang, and Y.-L. Chang, "Accelerating the kalman filter on a gpu," in *Parallel and Distributed Systems (ICPADS), 2011 IEEE 17th International Conference on*. IEEE, 2011, pp. 1016–1020.
- [5] O. Rosén and A. Medvedev, "Efficient parallel implementation of a kalman filter for single output systems on multicore computational platforms," in *Decision and Control and European Control Conference (CDC-ECC), 2011 50th IEEE Conference on*. IEEE, 2011, pp. 3178–3183.
- [6] C. Liu, "cuhmm: a cuda implementation of hidden markov model training and classification," *The Chronicle of Higher Education*, 2009.
- [7] J. Nielsen and A. Sand, "Algorithms for a parallel implementation of hidden markov models with a small state space," in *Parallel and Distributed Processing Workshops and Phd Forum (IPDPSW), 2011 IEEE International Symposium on*. IEEE, 2011, pp. 452–459.
- [8] M. Harris *et al.*, "Optimizing parallel reduction in cuda," *NVIDIA Developer Technology*, vol. 2, no. 4, 2007.
- [9] C. Nvidia, "Compute unified device architecture programming guide," 2007.
- [10] M. Harris, "Optimizing cuda," *SC07: High Performance Computing With CUDA*, 2007.
- [11] S. Johansen, "Statistical analysis of cointegration vectors," *Journal of economic dynamics and control*, vol. 12, no. 2, pp. 231–254, 1988.
- [12] F. Spitzer, *Principles of random walk*. Springer Science & Business Media, 2013, vol. 34.